

Unsupervised Part of Speech Induction Using Paradigmatic Representations of Word Context

Mehmet Ali Yatbaz*
Koç University

Enis Sert
Koç University

Deniz Yuret
Koç University

We investigate paradigmatic representations of word context in the domain of unsupervised part of speech induction. Paradigmatic representations of word context are based on potential substitutes of a word in contrast to syntagmatic representations based on properties of neighboring words. We demonstrate paradigmatic representations within two frameworks: (1) context clustering and (2) co-occurrence modeling. In context clustering we cluster word contexts based on the potential substitutes and they reveal a grouping that largely match the traditional part of speech boundaries. In co-occurrence modelling we construct a Euclidean embedding that models the co-occurrence of word types and their contexts. Clustering the points that correspond to word types in the Euclidean embedding gives state-of-the-art results in unsupervised part of speech induction, including 80% many-to-one accuracy on the Penn Treebank and significant improvements on 17 out of 19 corpora in 15 languages.


1. Introduction

Grammar rules apply not to individual words (e.g. dog, eat) but to syntactic categories of words (e.g. noun, verb). Thus constructing syntactic categories (also known as lexical or part of speech categories) is one of the fundamental problems in language acquisition.

Syntactic categories represent groups of words that can be substituted for one another without altering the grammaticality of a sentence. Linguists identify syntactic categories based on semantic, syntactic, and morphological properties of words. There is also evidence that children use prosodic and phonological features to bootstrap syntactic category acquisition (Ambridge and Lieven 2011). However there is as yet no satisfactory computational model that match human performance. Thus identifying the best set of features and best learning algorithms for syntactic category acquisition is still an open problem.

Relationships between linguistic units can be classified into two types: syntagmatic (concerning positioning), and paradigmatic (concerning substitution). Syntagmatic relations determine which units can combine to create larger groups and paradigmatic relations determine which units can be substituted for one another. Figure 1 illustrates the paradigmatic vs syntagmatic axes for words in a simple sentence and their possible substitutes.

* Artificial Intelligence Laboratory, Koç University, 34450 Sarıyer, İstanbul, Turkey. E-mail: myatbaz@ku.edu.tr, esert@ku.edu.tr, dyuret@ku.edu.tr

**Figure 1**

Syntagmatic vs. paradigmatic axes for words in a simple sentence (Chandler 2007).

In this study, we represent the paradigmatic axis directly by building *substitute vectors* for each word position in the text. The dimensions of a substitute vector represent words in the vocabulary, and the magnitudes represent the probability of occurrence in the given position. Note that the substitute vector for a word position (e.g. the second word in Fig. 1) is a function of the context only (i.e. “the ___ cried”), and does not depend on the word that does actually appear there (i.e. “man”). Thus substitute vectors represent *individual word contexts*, not word types. We refer to the use of features based on substitute vectors as *paradigmatic representations of word context*.

The two examples below illustrate the advantage of paradigmatic representations in uncovering similarities where no overt similarity that can be captured by a syntagmatic representation exists. The word “board” from the first sentence and the word “council” from the second sentence have no common neighbors except the determiner “the”. The paradigmatic representation captures the similarity of these words by suggesting the same top substitutes for both (the numbers in parentheses give substitute probabilities):

(1) “*Pierre Vinken, 61 years old, will join the **board** as a nonexecutive director Nov. 29.*”

board: board (.4288), company (.2584), firm (.2024), bank (.0731), ...

(2) “*... and hold only 25 % of the seats on the **council** .*”

council: board (.6591), company (.0795), firm (.0542), bank (.0154), ...

The high probability substitutes reflect both semantic and syntactic properties of the context. Top substitutes for “board” and “council” are not only nouns, but specifically nouns compatible with the semantic context. Top substitutes for the word “the” in the first example consist of words that can act as determiners: its (.9011), the (.0981), a (.0006), ...

There are two concerns inherent in all distributional methods: (i) words that are generally substitutable like “the” and “its” are placed in separate categories (DT and PRP\$) by the gold standard, (ii) words that are generally not substitutable like “do” and “put” are placed in the same category (VB). Freudenthal et al. (2005) point out that categories with unsubstitutable words fail the standard linguistic definition of a syntactic category and children do not seem to make errors of substituting such words in utterances (e.g. “*What do you want?*” vs. **“What put you want?”*). Whether gold standard part of speech tags or distributional categories are better suited to applications like parsing or machine translation can be best decided using extrinsic evaluation. In this study we evaluate our results by comparing them to gold standard part of speech tags and leave the extrinsic evaluation of the induced tags for future work.

Our preliminary experiments on a subsection of 1M word Penn Treebank Wall Street Journal corpus (PTB) (Marcus et al. 1999) indicate that using context information alone without the identity or the features of the target word (e.g. using dimensionality reduction and clustering on substitute vectors) has limited success and modeling the co-occurrence of word and context types is essential for inducing syntactic categories. In order to do so, we combine paradigmatic representations of word context with features of co-occurring words within the co-occurrence data embedding (CODE) framework (Globerson et al. 2007; Maron, Lamar, and Bienenstock 2010). The resulting embeddings for word types are split into 45 clusters using k-means and the clusters are compared to the 45 gold tags in the PTB. We obtain many-to-one accuracies up to .7680 using only distributional information (the identity of the word and a representation of its context) and .8023 by adding morphological and orthographic features of the words improving the state-of-the-art in unsupervised part of speech induction performance on the PTB. We extend the experiments to 19 corpora in 15 languages and achieve state-of-the-art many-to-one (MTO) scores on 17 of them.

The next section gives a detailed review of related work. Section 3 describes the construction of the substitute vectors and applies various distance metrics, dimensionality reduction and clustering methods on substitute vectors of the PTB. Section 4 describes co-occurrence data embedding (i.e. the learning algorithm used in our later experiments), introduces possible usage scenarios of substitute vectors and determines the best setup on the PTB. Section 5 extends experiments to 15 languages and compares our results with the best published results. Section 6 gives a brief error analysis and Section 7 summarizes our contributions. All the data and the code to replicate the results given in this paper is available from the authors' website at xxx.xxx.xxx.

2. Related Work

There are several good reviews of algorithms for unsupervised part of speech induction (Christodoulopoulos, Goldwater, and Steedman 2010; Gao and Johnson 2008) and models of syntactic category acquisition (Ambridge and Lieven 2011).

This work is to be distinguished from supervised part of speech disambiguation systems, which use labeled training data (Toutanova et al. 2003), unsupervised disambiguation systems, which use a dictionary of possible tags for each word (Yatbaz and Yuret 2010), or prototype driven systems which use a small set of prototypes for each class (Haghighi and Klein 2006). The problem of induction is important for studying under-resourced languages that lack labeled corpora and high quality dictionaries. It is also essential in modeling child language acquisition because every child manages to induce syntactic categories without access to labeled sentences, labeled prototypes, or dictionary constraints.

Models of unsupervised part of speech induction fall into two broad groups based on the information they utilize. Distributional models only use word types and their context statistics. Word-feature models incorporate additional morphological and orthographic features.

2.1 Distributional models

Distributional models can be further categorized into three subgroups based on the learning algorithm. The first subgroup represents each word type with its context vector and clusters these vectors accordingly (Schütze 1995). Work in modeling child syntactic category acquisition has generally followed this clustering approach (Redington, Crater, and Finch 1998; Mintz 2003). The second subgroup consists of probabilistic models based on the Hidden Markov Model (HMM) framework (Brown et al. 1992). A third group of algorithms constructs a low

dimensional representation of the data that represents the empirical co-occurrence statistics of word types (Globerson et al. 2007), which is covered in more detail in Section 4.

Clustering. Clustering based methods represent context using neighboring words, typically a single word on the left and a single word on the right called a “frame” (e.g., **the dog is; the cat is**). They cluster word types rather than word tokens based on the frames they occupy thus employing one-tag-per-word assumption from the beginning (with the exception of some methods in (Schütze 1995)). They may suffer from data sparsity caused by infrequent words and infrequent contexts. The solutions suggested either restrict the set of words and set of contexts to be clustered to the most frequently observed, or use dimensionality reduction. Redington et al. (1998) define context similarity based on the number of common frames bypassing the data sparsity problem but achieve mediocre results. Mintz (2003) only uses the most frequent 45 frames and Biemann (2006) clusters the most frequent 10,000 words using contexts formed from the most frequent 150-200 words. Schütze (1995) and Lamar et al. (2010) employ SVD to enhance similarity between less frequently observed words and contexts. Lamar et al. (2010) represent each context by the currently assigned left and right tag (which eliminates data sparsity) and cluster word types using a soft k-means style iterative algorithm. They report the best clustering result to date of .708 many-to-one accuracy on the PTB.

HMMs. The prototypical bitag HMM model maximizes the likelihood of the corpus $w_1 \dots w_n$ expressed as $P(w_1|c_1) \prod_{i=2}^n P(w_i|c_i)P(c_i|c_{i-1})$ where w_i are the word tokens and c_i are their (hidden) tags. One problem with such a model is its tendency to distribute probabilities equally and the resulting inability to model highly skewed word-tag distributions observed in hand-labeled data (Johnson 2007). To favor sparse word-tag distributions one can enforce a strict one-tag-per-word solution (Brown et al. 1992; Clark 2003), use sparse priors in a Bayesian setting (Goldwater and Griffiths 2007; Johnson 2007), or use posterior regularization (Ganchev et al. 2010). Each of these techniques provide significant improvements over the standard HMM model: for example Gao and Johnson (2008) show that sparse priors can gain from 4% (.62 to .66 on the PTB) in cross-validated many-to-one accuracy. However Christodoulopoulos et al. (2010) show that the older one-tag-per-word models such as (Brown et al. 1992) outperform the more sophisticated sparse prior and posterior regularization methods both in speed and accuracy (the Brown model gets .68 many-to-one accuracy on the PTB). Given that 93.69% of the word occurrences in human labeled data are tagged with their most frequent part of speech (Toutanova et al. 2003), this is probably not surprising; one-tag-per-word is a fairly good first approximation for induction.

2.2 Word-feature models

One problem with the algorithms in the previous section is the poverty of their input features. Of the syntactic, semantic, and morphological information linguists claim underlie syntactic categories, context vectors or bitag HMMs only represent limited syntactic information in their input. Experiments incorporating morphological and orthographic features into HMM based models demonstrate significant improvements. (Clark 2003; Berg-Kirkpatrick and Klein 2010; Blunsom and Cohn 2011) incorporate similar orthographic features and report improvements of 3, 7, and 10% respectively over the baseline Brown model. Christodoulopoulos et al. (2010) use prototype based features as described in (Haghighi and Klein 2006) with automatically induced prototypes and report an 8% improvement over the baseline Brown model. Abend et al. (2010) train a prototype-driven model with morphological features by first clustering the high frequency words as the landmarks and then assigning the remaining words to the landmark clusters. Christodoulopoulos et al. (2011) define a type-based Bayesian multinomial mixture

model in which each word instance is generated from the corresponding word type mixture component and word contexts are represented as features. They achieve a .728 MTO score by extending their model with additional morphological and alignment features gathered from parallel corpora. To our knowledge, nobody has yet tried to incorporate phonological or prosodic features in a computational model for syntactic category acquisition.

2.3 Paradigmatic representations

Yatbaz et al. (2012) explore the paradigmatic representation of word contexts by modeling the co-occurrence of words and their substitutes within the CODE framework. Their experiments on the PTB shows that paradigmatic representation improves the state-of-the-art MTO and V-measure (VM) accuracies of both distributional models and models with additional word features. This paper builds on that preliminary work by (1) exploring clustering of substitute vectors, (2) improving the model for using additional features, and (3) experimenting with additional languages.

Sahlgren (2006) gives a detailed analysis of paradigmatic and syntagmatic relations in the context of word-space models used to represent word meaning. Sahlgren’s paradigmatic model represents word types using co-occurrence counts of their frequent neighbors, in contrast to his syntagmatic model that represents word types using counts of contexts (documents, sentences) they occur in. Our substitute vectors do not represent word types at all, but *contexts of word tokens* using probabilities of likely substitutes. Sahlgren finds that in word-spaces built by frequent neighbor vectors, more nearest neighbors share the same part of speech compared to word-spaces built by context vectors.

Similarly, Schütze and Pedersen (1993) define the words that frequently co-occur together as the *syntagmatic associates* and words that have similar left and right neighbors as the *paradigmatic parallels*. Turney and Pantel (2010) give a broad overview of syntagmatic approaches and their applications within the Vector Space Modeling framework. We find that representing the paradigmatic axis more directly using substitute vectors rather than frequent neighbors improves part of speech induction.

Our paradigmatic representation is also related to the second order co-occurrences used in (Schütze 1995). Schütze concatenates the left and right context vectors for the target word type with the left context vector of the right neighbor and the right context vector of the left neighbor. The vectors from the neighbors include potential substitutes. Our method improves on his foundation by using a 4-gram language model rather than bigram statistics, using the whole 78,498 word vocabulary rather than the most frequent 250 words. More importantly, rather than simply concatenating vectors that represent the target word with vectors that represent the context we use a co-occurrence modeling algorithm.

2.4 Evaluation

We report many-to-one and V-measure scores for our experiments as suggested in (Christodoulopoulos, Goldwater, and Steedman 2010). The many-to-one (MTO) evaluation maps each cluster to its most frequent gold tag and reports the percentage of correctly tagged instances. The MTO score naturally gets higher with increasing number of clusters but it is an intuitive metric when comparing results with the same number of clusters. The V-measure (VM) (Rosenberg and Hirschberg 2007) is an information theoretic metric that reports the harmonic mean of homogeneity (each cluster should contain only instances of a single class) and completeness (all instances of a class should be members of the same cluster). In Section 6 we argue that homogeneity is perhaps more important in part of speech induction and suggest MTO with a fixed number of clusters as a more intuitive metric.

3. Substitute Vectors

In this section we explore the usage of substitute vectors within the context clustering framework by comparing similarity metrics, dimensionality reduction techniques and clustering methods. The first subsection describes the computation of substitute vectors using a statistical language model. Section 3.2 gives a detailed comparison of similarity metrics in the high dimensional substitute vector space. Section 3.3 analyzes the application of dimensionality reduction algorithms to substitute vectors. Section 3.4 presents a comparison of clustering methods on the PTB.

3.1 Computation of Substitute Vectors

In this study, we predict the syntactic category of a word in a given context based on its substitute vector. The dimensions of the substitute vector represent words in the vocabulary, and the entries in the substitute vector represent the probability of those words being used in the given context. Note that the substitute vector is a function of the context only and is indifferent to the target word.

It is best to use both the left and the right context when estimating the probabilities for potential lexical substitutes. For example, in “*He lived in San Francisco suburbs.*”, the token *San* would be difficult to guess from the left context but it is almost certain looking at the right context. We define c_w as the $2n - 1$ word window centered around the target word position: $w_{-n+1} \dots w_0 \dots w_{n-1}$ ($n = 4$ is the n -gram order). The probability of a substitute word w in a given context c_w can be estimated as:

$$P(w_0 = w | c_w) \propto P(w_{-n+1} \dots w_0 \dots w_{n-1}) \quad (1)$$

$$= P(w_{-n+1})P(w_{-n+2} | w_{-n+1}) \dots P(w_{n-1} | w_{-n+1}^{n-2}) \quad (2)$$

$$\approx P(w_0 | w_{-n+1}^{-1})P(w_1 | w_{-n+2}^0) \dots P(w_{n-1} | w_0^{n-2}) \quad (3)$$

where w_i^j represents the sequence of words $w_i w_{i+1} \dots w_j$. In Equation 1, $P(w | c_w)$ is proportional to $P(w_{-n+1} \dots w_0 \dots w_{n-1})$ because the words of the context are fixed. Terms without w_0 are identical for each substitute in Equation 2 therefore they have been dropped in Equation 3. Finally, because of the Markov property of n -gram language model, only the closest $n - 1$ words are used in the experiments.

Near the sentence boundaries the appropriate terms were truncated in Equation 3. Specifically, at the beginning of the sentence shorter n -gram contexts were used and at the end of the sentence terms beyond the end-of-sentence token were dropped.

To compute substitute probabilities we trained a language model using approximately 126 million tokens of Wall Street Journal data (1987-1994) extracted from CSR-III Text (Graff, Rosenfeld, and Paul 1995) (excluding sections of the PTB). We used SRILM (Stolcke 2002) to build a 4-gram language model with Kneser-Ney discounting. Words that were observed less than 500 times in the LM training data were replaced by UNK tags, which gave us a vocabulary size of 12,672. The first 24,020 tokens of the Penn Treebank Wall Street Journal Section 00 (PTB24K) was used as the test corpus to be induced. The corpus size was kept small in order to efficiently compute full distance matrices. Substitution probabilities for 12,672 vocabulary words were computed at each of the 24,020 positions. The perplexity of the 4-gram language model on the test corpus was 55.4 which is quite low due to using a small vocabulary and in-

domain data. The treebank uses 45 part of speech tags which is the set we used as the gold standard for comparison in our experiments.

3.2 Distance Metrics

We represent each context with a sparse high dimensional probability vector called the substitute vector as described in the previous section. In this section we compare various distance metrics in this high dimensional space with the goal of discovering one that will judge vectors that belong to the same syntactic category similar and vectors that belong to different syntactic categories distant. The distance metrics we have considered are listed in Table 1.

Table 1

Similarity metrics. JS is the Jensen-Shannon divergence and KL2 is a symmetric implementation of Kullback-Leibler divergence. Bold lower case letters represent vectors.

Cosine (p , q)	=	$\langle \mathbf{p}, \mathbf{q} \rangle / (\ \mathbf{p}\ _2 \ \mathbf{q}\ _2)$
Euclid (p , q)	=	$\ \mathbf{p} - \mathbf{q}\ _2$
Manhattan (p , q)	=	$\ \mathbf{p} - \mathbf{q}\ _1$
Maximum (p , q)	=	$\ \mathbf{p} - \mathbf{q}\ _\infty$
KL2 (p , q)	=	$\sum_i p_i \ln(p_i/q_i) + q_i \ln(q_i/p_i)$
JS (p , q)	=	$\sum_i p_i \ln(p_i/m_i) + q_i \ln(q_i/m_i)$ where $m_i = 0.5(p_i + q_i)$

To judge the merit of each distance metric we obtained supervised baseline scores using leave-one-out cross validation and the weighted k-nearest-neighbor algorithm¹ on the gold tags of the PTB24K. The results are listed in Table 2 sorted by score.

Table 2

Supervised baseline scores with different distance metrics. Log-metric indicates that metric applied to the log of the probability vectors.

Metric	Accuracy
KL2	.6889
Manhattan	.6865
Jensen	.6801
Cosine	.6706
Maximum	.6663
Euclid	.6255
log-Maximum	.5361
log-Cosine	.4847
log-Euclid	.4038
log-Manhattan	.3729

The entries with the log-prefix indicate a metric applied to the log of the probability vectors. Distance metrics on log probability vectors performed poorly compared to their regular counterparts indicating that differences in low probability words are relatively unimportant and high probability substitutes determine syntactic category. The surprisingly good result achieved by

¹ Neighbors were weighted using $1/\text{distance}$, $k = 30$ was chosen empirically.

the simple Maximum metric (which identifies the dimension with the largest difference between two vectors) also support this conclusion. The maximum score of .68 can be taken as a rough upper bound for an unsupervised learner using this space on the PTB24K corpus because 32% of the instances are assigned to the wrong part of speech by the majority of their closest neighbors. We will discuss alternative ways to push this upper bound higher by enforcing a one-tag-per-word rule in Section 3.4 and considering the context together with the co-occurring word type in Section 4.

3.3 Dimensionality Reduction

Using high dimensional vectors is problematic with many learning algorithms because of the computational cost and the curse of dimensionality. In this section we investigate if there is a low dimensional representation of the substitute vectors which still preserve the neighborhood information necessary to learn syntactic categories. We first briefly describe then report experimental results on principal components analysis (PCA), Isomap (Tenenbaum, Silva, and Langford 2000), locally linear embedding (LLE) (Roweis and Saul 2000), and Laplacian eigenmaps (Belkin and Niyogi 2003).

Each dimensionality reduction algorithm tries to preserve certain aspects of the original vectors. PCA is a linear method that minimizes reconstruction error. Isomap tries to preserve distances as measured along a low dimensional submanifold assuming the input vectors were sampled from the neighborhood of such a manifold. LLE most faithfully preserves the local linear structure of nearby input vectors. Laplacian eigenmaps most faithfully preserve proximity relations, mapping nearby inputs to nearby outputs.

We wanted to see how accuracy (based on the k-nearest-neighbor supervised baseline as in the previous section) changes based on the number of dimensions for each dimensionality reduction algorithm. For algorithms that require a distance matrix rather than raw input vectors we used the Jensen-Shannon divergence judged best by the experiments of the previous section. For graph based methods we built neighborhood graphs using 100 nearest neighbors. The low dimensional output vectors were compared using the cosine distance metric for the supervised k-nearest-neighbor algorithm. Figure 2 plots supervised baseline accuracy vs. number of dimensions for each algorithm.




Figure 2
Supervised knn baselines for the four dimensionality reduction algorithms.

The graph based algorithms (Isomap, LLE, and Laplacian eigenmaps) all outperform PCA. They stay within 5% of their peak accuracy with as few as 16 dimensions. In fact Laplacian eigenmaps outperform the baseline with the original 12,672 dimensional vectors (68.95%) when allowed to retain more than about 250 dimensions. Spectral clustering uses the same transformation as the Laplacian eigenmaps algorithm and we compare its performance to other clustering algorithms in the next section.

3.4 Clustering

We compared three clustering algorithms applied to the original substitute vectors using many-to-one accuracy on the PTB24K. Hierarchical agglomerative clustering with complete linkage (HAC) starts with each instance in its own cluster and iteratively combines the two closest groups (measured by their most distant points) at each step (Manning, Raghavan, and Schütze 2008). K-medoids minimizes sum of pairwise distances between each data-point to the exemplar at the center of its cluster (Kaufman and Rousseeuw 2005). Spectral clustering² uses the eigenvalues of the graph Laplacian $L = D^{-1/2}WD^{-1/2}$ to reduce the number of dimensions (similar to Laplacian eigenmaps) and uses simple k-means clustering on the resulting representation (Ng, Jordan, and Weiss 2002). All three algorithms accept the distance matrix based on the KL2 distance (see Section 3.2) as input.




Figure 3

Many-to-one score for three clustering algorithms on the 45-tag 24K word corpus.

Figure 3 plots the many-to-one score versus number of clusters for the three algorithms on the PTB24K. The many-to-one score naturally increases as we approach the one cluster per word limit, however we find the evolution of the curves informative. At the high end (more than 2000 clusters) HAC performs best with its conservative clusters, but its performance degrades fast as we reduce the number of clusters because it cannot reverse the accumulating mistakes. At the low end (less than 16 clusters) k-medoids and spectral have similar performance. However for the region of interest (between 16 to 2000 clusters) spectral clustering is clearly superior with .5841 MTO accuracy at 45 clusters.

To demonstrate that the clustering algorithm can scale, we performed spectral clustering on the full PTB. We calculated the substitute vectors of each position in the PTB by using an

² We used the implementation in (Chen et al. 2011) with a symmetric sparse affinity matrix of 550 nearest neighbors.

efficient algorithm described in Section 4.2³. We used the Manhattan distance⁴ to compute the vector similarities and achieve .5473 MTO accuracy.

We noted that the 45 cluster spectral clustering result assigned many more tags to each word than the gold standard. To increase the sparsity of the solution we enforced a one-tag-per-word restriction by re-tagging each word with its most frequent tag in the original assignment (we break ties randomly). One-tag-per-word re-tagging improves the MTO accuracy from .5473 to .6834.

The re-tagging experiment shows that the one-tag-per-word restriction improves the model performance in terms of the MTO accuracy by explicitly incorporating the word type information into the model. In the next section we suggest an alternative method that models the co-occurrence of word types and their contexts as represented by substitutes and improves the results on the PTB.

4. Co-occurrence Modeling

In this section we combine the paradigmatic representation of the word context with the identity and features of the target word for part of speech induction using a co-occurrence modeling framework. Our preliminary experiments in Section 3.4 indicated that using the context information alone (e.g. clustering substitute vectors) without the target word identity and features had limited success. Moreover incorporating word identities (i.e. the one-tag-per-word constraint) even in an ad-hoc manner by re-tagging the clustering output significantly improves the MTO accuracy. It seems that the co-occurrence of a target word with a particular type of context may be a better predictor of the syntactic category.

Section 4.1 reviews the unsupervised methods we use to model co-occurrence statistics: the Co-occurrence Data Embedding (CODE) (Globerson et al. 2007) method and its spherical extension (S-CODE) introduced by (Maron, Lamar, and Bienenstock 2010). Section 4.2 describes the general experimental settings. The S-CODE algorithm works with discrete inputs. The substitute vectors as described in Section 3 are high dimensional and continuous. We experimented with two approaches to use substitute vectors in a discrete setting. Section 4.3 presents an algorithm that partitions the high dimensional space of substitute vectors into small neighborhoods and uses the partition id as a discrete context representation. Section 4.4 presents an even simpler model which pairs each word with random substitutes sampled from the substitute vector. Section 4.5 replicates the bigram based S-CODE results from (Maron, Lamar, and Bienenstock 2010) as a direct comparison of paradigmatic vs syntagmatic representations of word context. When the left-word – right-word pairs used in the bigram model are replaced with word – partition-id or word – substitute pairs we see significant gains in accuracy. These results support our running hypothesis that paradigmatic features, i.e. potential substitutes of a word, are better determiners of syntactic category compared to left and right neighbors (syntagmatic features). Section 4.6 explores morphological and orthographic features as additional sources of information and its results improve the state-of-the-art in the field of unsupervised part of speech induction.

³ The algorithm in Section 4.2 only calculates the top 100 substitutes and their unnormalized probabilities which results in sparse substitute vectors. As a result the new model is computationally more efficient than the one in this section and it is more appropriate for large datasets like the PTB.

⁴ KL2 is undefined on sparse vectors, therefore we use Manhattan whose performance is close to KL2 on both the PTB24K and PTB.

4.1 The CODE Model

Let X and Y be two categorical variables with finite cardinalities $|X|$ and $|Y|$. We observe a set of pairs $\{x_i, y_i\}_{i=1}^n$ drawn IID from the joint distribution of X and Y . The basic idea behind CODE and related methods is to represent (embed) each value of X and each value of Y as points in a common low dimensional Euclidean space \mathbf{R}^d such that values that frequently co-occur lie close to each other. There are several ways to formalize the relationship between the distances and co-occurrence statistics, in this paper we use the following:

$$p(x, y) = \frac{1}{Z} \bar{p}(x) \bar{p}(y) e^{-d_{x,y}^2} \quad (4)$$

where $d_{x,y}^2$ is the squared distance between the embeddings of x and y , $\bar{p}(x)$ and $\bar{p}(y)$ are empirical probabilities, and $Z = \sum_{x,y} \bar{p}(x) \bar{p}(y) e^{-d_{x,y}^2}$ is a normalization term. If we use the notation ϕ_x for the point corresponding to x and ψ_y for the point corresponding to y then $d_{x,y}^2 = \|\phi_x - \psi_y\|^2$. The log-likelihood of a given embedding $\ell(\phi, \psi)$ can be expressed as:

$$\begin{aligned} \ell(\phi, \psi) &= \sum_{x,y} \bar{p}(x, y) \log p(x, y) \\ &= \sum_{x,y} \bar{p}(x, y) (-\log Z + \log \bar{p}(x) \bar{p}(y) - d_{x,y}^2) \\ &= -\log Z + \text{const} - \sum_{x,y} \bar{p}(x, y) d_{x,y}^2 \end{aligned} \quad (5)$$

The likelihood is not convex in ϕ and ψ . We use gradient ascent to find an approximate solution for a set of ϕ_x, ψ_y that maximize the likelihood. The gradient of the $d_{x,y}^2$ term pulls neighbors closer in proportion to the empirical joint probability:

$$\frac{\partial}{\partial \phi_x} \sum_{x,y} -\bar{p}(x, y) d_{x,y}^2 = \sum_y 2\bar{p}(x, y) (\psi_y - \phi_x) \quad (6)$$

The gradient of the Z term pushes neighbors apart in proportion to the estimated joint probability:

$$\frac{\partial}{\partial \phi_x} (-\log Z) = \sum_y 2p(x, y) (\phi_x - \psi_y) \quad (7)$$

Thus the net effect is to pull pairs together if their estimated probability is less than the empirical probability and to push them apart otherwise. The gradients with respect to ψ_y are similar.

S-CODE (Maron, Lamar, and Bienenstock 2010) additionally restricts all ϕ_x and ψ_y to lie on the unit sphere. With this restriction, Z stays around a fixed value during gradient ascent. This allows S-CODE to substitute an approximate constant \tilde{Z} in gradient calculations for the real Z for computational efficiency. In our experiments, we used S-CODE with its sampling based stochastic gradient ascent algorithm and smoothly decreasing learning rate.

4.2 Experimental Settings

To make a meaningful comparison on the PTB we ran all the experiments using the following default settings (unless otherwise stated): (i) each word was kept with its original capitalization,

Table 3

Summary of results in terms of the MTO and VM scores. Standard errors are given in parentheses when available. Starred entries have been reported in the review paper (Christodoulopoulos, Goldwater, and Steedman 2010). Distributional models use only the identity of the target word and its context. The models on the right incorporate orthographic and morphological features.

Distributional Models	MTO	VM	Models with Additional Features	MTO	VM
Lamar et al. (2010)	.708	-	Clark (2003)*	.712	.655
Brown et al. (1992)*	.678	.630	Christodoulopoulos et al. (2011)	.728	.661
Goldwater et al. (2007)	.632	.562	Berg-Kirkpatrick et al. (2010)	.755	-
Ganchev et al. (2010)*	.625	.548	Christodoulopoulos et al. (2010)	.761	.688
Maron et al. (2010)	.688 (.0016)	-	Blunsom and Cohn (2011)	.775	.697
Bigrams (Sec. 4.5)	.7314 (.0096)	.6558 (.0052)	Substitutes and Features (Sec. 4.6)	.8004 (.0075)	.7160 (.0044)
Partitions (Sec. 4.3)	.7554 (.0055)	.6703 (.0037)			
Substitutes (Sec. 4.4)	.7680 (.0038)	.6822 (.0029)			

(ii) the learning rate parameters were set to $\varphi_0 = 50$, $\eta_0 = 0.2$ for faster convergence in log likelihood, (iii) the number of S-CODE iterations were set to 50 million, (iv) the S-CODE dimensions and Z were set to 25 and 0.166, respectively, (v) a modified k-means algorithm with smart initialization was used (Arthur and Vassilvitskii 2007), and (vi) the number of k-means restarts were set to 128 to improve clustering and reduce variance.

Section 3.2 shows that low probability substitutes are relatively unimportant thus for computational efficiency only the top 100 substitutes and their unnormalized probabilities were computed for each positions in the PTB (1,173,766 tokens, 49,206 types) using the FASTSUBS algorithm (Yuret 2012)⁵. The probability vectors for each position were normalized to add up to 1.0 giving us the final substitute vectors used in the rest of this study. We set the vocabulary threshold to 20 which increases the vocabulary size to 78,498.

Each experiment was repeated 10 times with different random seeds and the results are reported with standard errors in parentheses or error bars in graphs. Table 3 summarizes all the results reported in this section and the ones we cite from the literature.

4.3 Random Partitions

To obtain a discrete representation of the context, the random-partitions algorithm first designates a random subset of substitute vectors as centroids to partition the space, and then associates each context with the partition defined by the closest centroid in cosine distance. Each partition thus defined gets a unique id, and word (X) – partition-id (Y) pairs are given to S-CODE as input. The algorithm uses stochastic gradient ascent to find the ϕ_x, ψ_y embeddings for word and partition-id in these pairs on a single 25-dimensional sphere. The algorithm cycles through the data until we get approximately 50 million updates. The resulting ϕ_x vectors are clustered using an instance weighted k-means algorithm and the resulting groups are compared to the correct part of speech tags. Using default settings with 64K random partitions the many-to-one accuracy is .7554 (.0055) and the V-measure is .6703 (.0037).

To analyze the sensitivity of this result to our specific parameter settings we ran a number of experiments where each parameter was varied over a range of values.

Figure 4 gives results where the number of initial random partitions is varied over a large range and shows the results to be fairly stable across two orders of magnitude.

⁵ The substitutes with unnormalized log probabilities can be downloaded from <http://goo.gl/jzKH0>.




Figure 1 is a plot showing the mass-to-light ratio m_{20} (y-axis, ranging from 0.7 to 0.8) versus the s-code Z approximation (x-axis, logarithmic scale from 0.01 to 1). The data points are represented by black crosses with vertical error bars. A solid black line represents a fit to the data. The curve starts at approximately 0.71 for $s\text{-code } Z \approx 0.01$, rises sharply to a peak of about 0.765 at $s\text{-code } Z \approx 0.15$, and then gradually decreases to about 0.70 at $s\text{-code } Z \approx 1.5$.

Figure 6 shows that the constant \tilde{Z} approximation can be varied within two orders of magnitude without a significant performance drop in the many-to-one score. For uniformly distributed points on a 25 dimensional sphere, the expected $Z \approx 0.146$. In the experiments where

we tested we found the real Z always to be in the 0.140-0.170 range. When the constant \tilde{Z} estimate is too small the attraction in Eq. 6 dominates the repulsion in Eq. 7 and all points tend to converge to the same location. When \tilde{Z} is too high, it prevents meaningful clusters from coalescing.

We find the random partition algorithm to be fairly robust to different parameter settings and the resulting many-to-one score significantly better than the state-of-the-art distributional models.

4.4 Random Substitutes

Another way to use substitute vectors in a discrete setting is simply to sample individual substitute words from them according to the corresponding probabilities. The random-substitutes algorithm cycles through the test data and pairs each word with a random substitute picked from the pre-computed substitute vectors (see Section 3). We ran the random-substitutes algorithm to generate 76 million word (X) – random-substitute (Y) pairs (64 substitutes for each token) as input to S-CODE. Clustering the resulting ϕ_x vectors yields a many-to-one score of .7680 (.0038) and a V-measure of .6822 (.0029).

This result is close to the previous result by the random-partition algorithm, .7554 (.0055), demonstrating that two very different discrete representations of context based on paradigmatic features give consistent results. Figure 7 illustrates that the random-substitute result is fairly robust as long as the training algorithm can observe more than a few random substitutes per word.




Figure 7

MTO is not sensitive to the number of random substitutes sampled per word token.

4.5 Paradigmatic vs Syntagmatic Representations of Word Context

To get a direct comparison of the paradigmatic and syntagmatic context representations we feed the adjacent word pairs (bigrams) in the corpus into the S-CODE algorithm as X, Y samples (Maron, Lamar, and Bienenstock 2010) instead of pairing each word with a paradigmatic representation of its context. At the end each word w in the vocabulary ends up with two points on the sphere, a ϕ_w point representing the behavior of w as the left word of a bigram and a ψ_w point representing it as the right word. The two vectors for w are concatenated to create a 50-dimensional representation at the end. These 50-dimensional vectors are clustered using the k-means algorithm. Maron et al. (2010) report many-to-one scores of .6880 (.0016) for 45 clusters and .7150 (.0060) for 50 clusters (on the PTB). If only ϕ_w vectors are clustered without concatenation we found the performance drops significantly to about .62. Using our default

settings the bigram model achieves .7314 (.0096) MTO and .6558 (.0052) VM accuracies. Both results are significantly lower than the random partition and substitute MTO and VM accuracies.

4.6 Morphological and orthographic features

Clark (2003) demonstrates that using morphological and orthographic features significantly improves part of speech induction with an HMM based model. Section 2 describes a number of other approaches that show similar improvements. This section describes one way to integrate additional features to the random-substitute model.

In order to accommodate multiple feature types the CODE model needs to be extended to handle more than two variables. Globerson et al. (2007) suggest the following likelihood function:

$$\ell(\phi, \psi^{(1)}, \dots, \psi^{(K)}) = \sum_i^K w_i \sum_{x, y^{(i)}} \bar{p}(x, y^{(i)}) \log p(x, y^{(i)}) \quad (8)$$

where $Y^{(1)}, \dots, Y^{(K)}$ are K different variables whose empirical joint distributions with X , $\bar{p}(x, y^{(1)}) \dots \bar{p}(x, y^{(K)})$, are known. Eq. 8 then represents a set of CODE models $p(x, y^{(k)})$ where each $Y^{(k)}$ has an embedding $\psi_y^{(k)}$ but all models share the same ϕ_x embedding. The weights w_k reflect the relative importance of each $Y^{(k)}$ and all embeddings are mapped to unit-sphere.

We adopt this likelihood function, set all $w_k = 1$, let X represent a word, $Y^{(1)}$ represent a random substitute, and $Y^{(2)}, \dots, Y^{(K)}$ stand for morphological and orthographic features of the word thus each word is a $(K+1)$ -tuple, $(X, Y^{(1)}, \dots, Y^{(K)})$. With this setup, the training procedure needs to change little: instead of sampling a word – random-substitute pair, the word – random-substitute – features tuple is sampled and input to the gradient ascent algorithm. The gradient search algorithm updates the embeddings according to $p(x, y^{(i)})$ where $i = 1 \dots k$ and no updates are performed between $y^{(i)}$ s since they do not have any co-occurrence statistics and x is the shared variable.

Word tuples might have null values due to the unobserved features. For example, the word “car” has no morphological or orthographic features therefore all the elements of the tuple have null value except the word type (X) and the random-substitute ($Y^{(1)}$). We do not perform any pull or push updates on embeddings during the gradient search if the corresponding $y^{(k)}$ is null⁶.

The orthographic features we used are similar to the ones in (Berg-Kirkpatrick et al. 2010) with small modifications:

- Initial-Capital: this feature is generated for capitalized words with the exception of sentence initial words.
- Number: this feature is generated when the token starts with a digit.
- Contains-Hyphen: this feature is generated for lowercase words with an internal hyphen.
- Initial-Apostrophe: this feature is generated for tokens that start with an apostrophe.

⁶ X and $Y^{(1)}$ represents the word type and random-substitute therefore they are always observed.

We generated morphological features using the unsupervised algorithm Morfessor (Creutz and Lagus 2005). Morfessor was trained on the WSJ section of the Penn Treebank using default settings, and a perplexity threshold of 1. In our model, a word type consists of two parts: a stem and a suffix part. The suffix part is used as the morphological feature thus each word type has only one morphological feature⁷. The program induced 5575 suffix types that are present in a total of 19223 word types.

Using the training settings of the previous section, the addition of morphological and orthographic features increased the many-to-one score of the random-substitute model to .8004 (.0075) and V-measure to .7160 (.0044). Both these results improve the state-of-the-art in part of speech induction significantly as seen in Table 3.

4.7 Summary

We presented methods to determine the best usage of substitute vectors within the CODE framework and performed sensitivity analysis on the model parameters to not only show the robustness but also to decide the best configuration of S-CODE in modeling the co-occurrence of words with their contexts. To use the substitute vectors as co-occurrences we discretized the substitute vectors using two methods. The first method (Section 4.3) selected 64K random substitute vectors as the center of 64K partitions and then assigned rest of the substitute vectors to the closest partition. As a result each word represented as a word – partition-id pair which we input to S-CODE. The second method (Section 4.4) sampled random substitutes from the substitute vectors using the fact that the substitute vectors are probability distributions. We fed the word – random-substitute pairs into S-CODE. Both methods significantly outperform the syntagmatic bigram model (Section 4.5) on the PTB. Finally, Section 4.6 showed that adding morphological and orthographic features improved the accuracy and we achieved the state-of-the-art MTO and VM accuracies on the PTB.

In the next section we extend our experiments to include more languages to demonstrate the robustness of paradigmatic approach on languages with different characteristics.

5. Multilingual Experiments

We performed experiments with a range of languages and three different feature configurations to establish both the robustness of our model across languages and to observe the effects of different features. Following Christodoulopoulos et al. (2011), in addition to the PTB we extend our experiments to 8 languages from MULTEXT-East (Bulgarian, Czech, English, Estonian, Hungarian, Romanian, Slovene and Serbian) (Erjavec 2004) and 10 languages from the CoNLL-X shared task (Bulgarian, Czech, Danish, Dutch, German, Portuguese, Slovene, Spanish, Swedish and Turkish) (Buchholz and Marsi 2006). For all experiments, we use the best performing model of Section 4 (i.e. the random substitute model) with default settings. To perform meaningful comparisons with the previous work we train and evaluate our models on the training section of MULTEXT-East⁸ and CoNLL-X languages (Lee, Haghighi, and Barzilay 2010).

Section 5.1 details the language model and feature statistics of each language. Section 5.2 summarizes the results of our models for all of the languages in our corpora. In the rest of this section we refer to the MULTEXT-East and CoNLL-X corpora as the testing corpora and the language model training corpora as the training corpora.

⁷ We extracted the stem part by concatenating the splits until including the first “STM” labeled split and the suffix part by concatenating rest of the splits.

⁸ Languages of MULTEXT-East corpora do not tag the punctuation marks, thus we add an extra tag for punctuation to the tag-set of these languages.

5.1 Substitute Vectors and Features

Table 4

Summary of language model training and test corpora statistics for each language in the test set. Last two columns present the number of induced suffix parts and word types with these suffix parts after the morphological feature extraction.

	Language Model			Test set				
	Language	Source	Word Count	Word Count	Perplexity (ppl)	Unknown Word	Suffix Parts	Word Types with Suffix parts
WSJ	English	News	126,170,376	1,173,766	79.926	0.012	5575	19223
MULTEXT-East	Bulgarian	Wikipedia	32,511,616	101,173	655.202	.0565	609	4209
	Czech	Wikipedia	59,698,049	100,368	1,069.67	.0299	2787	12848
	English	News	126,170,376	118,424	265.246	.0288	1251	4783
	Estonian	Wikipedia	14,513,571	94,898	871.765	.0654	4448	13638
	Hungarian	Wikipedia	66,069,788	98,426	742.676	.0449	5423	15995
	Romanian	Wikipedia	35680870	118,328	666.855	.1074	2064	9445
	Slovene	Wikipedia	18,969,846	112,278	658.711	.0389	2093	11834
	Serbian	Wikipedia	17,129,679	108,809	804.962	.0580	2722	12476
CoNLL-X Shared Task	Bulgarian	Wikipedia	32,511,616	190,217	538.972	.0430	926	8225
	Czech	Wikipedia	59,698,049	1,249,408	1,233.95	.0250	12443	85673
	Danish	Wikipedia	35,863,945	94,386	351.24	.0393	3708	10897
	Dutch	Wikipedia	159,978,524	195,069	390.818	.0476	5250	13407
	German	Wikipedia	437,777,863	699,610	680.036	.0487	15219	45414
	Portuguese	Wikipedia	150,099,154	206,678	378.656	.0861	5033	15721
	Slovene	Wikipedia	18,969,846	28,750	663.053	.0414	1257	4781
	Spanish	Wikipedia	332,311,650	89,334	274.418	.0424	2648	9316
	Swedish	Wikipedia	32,004,538	191,467	1,233.95	.0250	2897	12725
	Turkish	Web	491,195,991	47,605	868.829	.0508	5651	14227

To calculate the top 100 substitutes of each position, we train a 4-gram language model with the corresponding training corpora of each language as described in Section 4.2. Table 4 presents statistics related to the language model training and testing corpora. For all languages except Serbian, English and Turkish, we train the language models by using the corresponding Wikipedia dump files⁹.

Serbian shares a common basis with Croatian and Bosnian therefore we trained 3 different language models using Wikipedia dump files of Serbian together with these two languages and measured the perplexities on the MULTEXT-East Serbian corpus. We chose the Croatian language model since it achieved the lowest perplexity score and unknown word ratio on the MULTEXT-East Serbian corpus.

To train the statistical language model of English, we use Wall Street Journal data (1987-1994) extracted from CSR-III Text (Graff, Rosenfeld, and Paul 1995) (excluding sections of the

⁹ Latest Wikipedia dump files are freely available at <http://dumps.wikimedia.org/> and the text in the dump files can be extracted using WP2TXT (<http://wp2txt.rubyforge.org/>)

PTB) and for the Turkish language modeling we use the web corpus collected from Turkish news and blog sites (Sak, Güngör, and Saraçlar 2008).

Language model training files vary across the languages in terms of quality and quantity. In order to reduce the unknown word ratio of resource poor languages and to standardize the process we set the vocabulary threshold to 2 for all languages except English. English has relatively low unknown word ratio therefore we set the threshold to 20 instead of 2.

We use the same set of orthographic features described in Section 4.6 except we add an “Only-Punctuation” feature to the languages of MULTEXT-East corpora. The “Only-Punctuation” feature is generated when a token only consists of punctuation characters.

Morphological features are extracted by the method described in Section 4.6 using the training sections of each language in MULTEXT-East and CoNLL-X corpora¹⁰. Language specific morphological feature statistics are summarized in Table 4.

Table 5

The MTO and VM scores on 19 corpora in 15 languages together with the number of types and tags in gold-set which equals to number of induced clusters in all languages. Best published results are from [‡](Blunsom and Cohn 2011), ^{*}(Christodoulopoulos, Goldwater, and Steedman 2011) and [†](Clark 2003). Bold results represent the best MTO and VM accuracies of the corresponding language. MULTEXT-East corpora do not tag the punctuation marks, thus we add an extra tag for punctuation and represent it with “+1”.

	Language	Types	Tags	Best Published	Syntagmatic Bigram	DIS	DIS+O	DIS+O+M
WSJ	English	49,190	45	.775 / .697 [‡]	.7314 / .6558	.7680 / .6822	.7830 / .7026	.8004 / .7160
MULTEXT-East	Bulgarian	16,352	12+1	.665 / .556 [*]	.6732 / .4119	.6883 / .5291	.7039 / .5496	.6754 / .5246
	Czech	19,115	12+1	.642 / .539 [*]	.6269 / .4586	.6781 / .4829	.6742 / .4854	.6977 / .5042
	English	9,773	12+1	.733 / .633 [*]	.7690 / .6131	.8229 / .6610	.8282 / .6719	.8343 / .6787
	Estonian	17,845	12+1	.644 / .533 [*]	.6089 / .4119	.6555 / .4437	.6634 / .4606	.6526 / .4418
	Hungarian	20,321	11+1	.682 / .548 [*]	.6181 / .4514	.6914 / .5046	.7052 / .5244	.7287 / .5444
	Romanian	15,189	14+1	.611 / .523 [*]	.6565 / .5202	.6469 / .5012	.6675 / .5269	.6488 / .5251
	Slovene	17,871	12+1	.679 / .567 [*]	.6772 / .5044	.6873 / .4845	.6892 / .4901	.6833 / .4941
	Serbian	18,095	12+1	.641 / .510 [†]	.6267 / .4510	.6240 / .4479	.6303 / .4554	.6368 / .4650
CoNLL-X Shared Task	Bulgarian	32,439	54	.704 / .596 [†]	.6972 / .5532	.7399 / .5824	.7391 / .5856	.7207 / .5673
	Czech	130,208	12	.701 [‡] / .484 [*]	.6944 / .5036	.6764 / .4867	.7149 / .5330	.6903 / .5227
	Danish	18,356	25	.761 [‡] / .591 [*]	.6757 / .5290	.7214 / .5559	.7520 / .5927	.7482 / .5958
	Dutch	28,393	13	.711 [‡] / .547	.6703 / .5205	.7014 / .5405	.7393 / .5980	.7228 / .5925
	German	72,326	54	.744 [*] / .630 [†]	.7525 / .6285	.7637 / .6314	.7735 / .6554	.7529 / .6403
	Portuguese	28,931	22	.785 [‡] / .639 [*]	.7031 / .5617	.7381 / .5770	.7907 / .6317	.7948 / .6405
	Slovene	7,128	29	.642 [*] / .539 [†]	.6384 / .4976	.6503 / .4925	.6555 / .5036	.6572 / .5023
	Spanish	16,458	47	.788 [‡] / .632 [*]	.7086 / .5844	.7492 / .6083	.7718 / .6372	.7627 / .6331
	Swedish	20,057	41	.682 / .589 [†]	.6721 / .5558	.6931 / .5654	.6946 / .5721	.6649 / .5613
	Turkish	17,563	30	.628 / .408 [*]	.6069 / .3551	.6228 / .3804	.6348 / .4109	.6500 / .4246

5.2 Results

For each language we report results of three models: (1) distributional (DIS), (2) distributional with orthographic features (DIS+O) and (3) distributional with both orthographic and morphological features (DIS+O+M). Similar to the settings used in Section 4.4, we use the

¹⁰ We don't use the language model corpora to extract morphological features.

25 dimensional sphere with 64 substitutes for all languages. For each language the number of induced clusters is set to the number of tags in the gold-set as presented in Table 5.

As a baseline model we chose the syntagmatic bigram version of S-CODE described in Section 4.5 which is a very strong baseline compared to the ones used in (Christodoulopoulos, Goldwater, and Steedman 2011). Table 5 summarizes the MTO and VM scores of our models together with the syntagmatic bigram baseline and the best published accuracies on each language corpus.

DIS significantly outperforms the syntagmatic bigram baseline in both MTO and VM scores on 14 languages. DIS+O+M has the state-of-the-art MTO and VM accuracy on the PTB. DIS+O and DIS+O+M achieve the highest MTO scores on all languages of MULTEXT-East corpora while scoring the highest VM accuracies on English and Romanian. On the CoNLL-X languages our models perform better than the best published MTO or VM accuracies on 10 languages.

6. Discussion

In this section we perform further analysis on the clustering output of our best model and indicate the possible reasons of comparably low VM scores. To illustrate how words are distributed in the induced clusters, we compare the output of our model with gold-tags of the PTB. We also discuss the effect of coarse gold-tag sets on our model performance.




Figure 8

Hinton diagram comparing most frequent tags and clusters. Area of each square is proportional to the joint probability of the given tag and cluster.

Figure 8 is the Hinton diagram of the PTB showing the relationship between the most frequent tags and clusters from the experiment in Section 4.6. In general the errors seem to be the lack of completeness (multiple large entries in a row), rather than lack of homogeneity (multiple large entries in a column). The algorithm tends to split large word classes into several clusters. Some examples are:

- Titles like Mr., Mrs., and Dr. are split from the rest of the proper nouns in cluster (39).

- Auxiliary verbs (10) and the verb “say” (22) have been split from the general verb clusters (12) and (7).
- Determiners “the” (40), “a” (15), and capitalized “The”, “A” (6) have been split into their own clusters.
- Prepositions “of” (19), and “by”, “at” (17) have been split from the general preposition cluster (8).

Nevertheless there are some homogeneity errors as well:

- The adjective cluster (5) also has some noun members probably due to the difficulty of separating noun-noun compounds from adjective modification.
- Cluster (6) contains capitalized words that span a number of categories.

Most closed-class items are cleanly separated into their own clusters as seen in the lower right hand corner of the diagram.

The completeness errors become more noticeable on languages with coarse tag-sets thus our models perform worse than the best published models on 6 of MULTEXT-East languages in terms of VM scores while achieving the state-of-the-art MTO scores on same languages as shown on Table 5. On CONLL-X languages the effect of completeness errors is less noticeable since all languages except Czech and Dutch have fine grained tag-sets.

The completeness errors are not surprising given that the words that have been split are not generally substitutable with the other members of their gold-tag set category. Thus it can be argued that metrics that emphasize homogeneity such as MTO are more appropriate in this context than metrics that average homogeneity and completeness such as VM as long as the number of clusters is controlled.

7. Contributions

Our main contributions can be summarized as follows:

- We introduced substitute vectors as paradigmatic representations of word context and demonstrated their use in unsupervised part of speech induction on 19 corpora in 15 languages.
- We demonstrated that using paradigmatic representations of word context and modeling co-occurrences of word and context types with the S-CODE learning framework give superior results when compared to a syntagmatic bigram model.
- We extended the S-CODE framework to incorporate morphological and orthographic features and improved the state-of-the-art many-to-one accuracy in unsupervised part of speech induction on 17 out of 19 corpora.
- All our code and data, including the substitute vectors for the PTB, MULTEXT-East and CoNLL-X shared task corpora are available at the authors’ website at `xxx.xxx.xxx`.

References

Abend, Omri, Roi Reichart, and Ari Rappoport. 2010. Improved unsupervised pos induction through prototype discovery. In *Proceedings of the 48th Annual Meeting of the Association for Computational*

- Linguistics*, ACL '10, pages 1298–1307, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Ambridge, B. and E.V.M. Lieven, 2011. *Child Language Acquisition: Contrasting Theoretical Approaches*, chapter 6.1. Cambridge University Press.
- Arthur, D. and S. Vassilvitskii. 2007. k-means++: The advantages of careful seeding. In *Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms*, pages 1027–1035. Society for Industrial and Applied Mathematics.
- Belkin, M. and P. Niyogi. 2003. Laplacian eigenmaps for dimensionality reduction and data representation. *Neural computation*, 15(6):1373–1396.
- Berg-Kirkpatrick, Taylor, Alexandre Bouchard-Côté, John DeNero, and Dan Klein. 2010. Painless unsupervised learning with features. In *Human Language Technologies: The 2010 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pages 582–590, Los Angeles, California, June. Association for Computational Linguistics.
- Berg-Kirkpatrick, Taylor and Dan Klein. 2010. Phylogenetic grammar induction. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics*, pages 1288–1297, Uppsala, Sweden, July. Association for Computational Linguistics.
- Biemann, C. 2006. Unsupervised part-of-speech tagging employing efficient graph clustering. In *Proceedings of the 21st International Conference on computational Linguistics and 44th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*, pages 7–12. Association for Computational Linguistics.
- Blunsom, Phil and Trevor Cohn. 2011. A hierarchical pitman-yor process hmm for unsupervised part of speech induction. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 865–874, Portland, Oregon, USA, June. Association for Computational Linguistics.
- Brown, Peter F., Peter V. deSouza, Robert L. Mercer, Vincent J. Della Pietra, and Jenifer C. Lai. 1992. Class-based n-gram models of natural language. *Comput. Linguist.*, 18:467–479, December.
- Buchholz, Sabine and Erwin Marsi. 2006. Conll-x shared task on multilingual dependency parsing. In *Proceedings of the Tenth Conference on Computational Natural Language Learning*, CoNLL-X '06, pages 149–164, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Chandler, D. 2007. *Semiotics: the basics*. The Basics Series. Routledge.
- Chen, WY, Y. Song, H. Bai, CJ Lin, and EY Chang. 2011. Parallel spectral clustering in distributed systems. *IEEE transactions on pattern analysis and machine intelligence*, 33(3):568–586.
- Christodoulopoulos, Christos, Sharon Goldwater, and Mark Steedman. 2010. Two decades of unsupervised pos induction: how far have we come? In *Proceedings of the 2010 Conference on Empirical Methods in Natural Language Processing*, EMNLP '10, pages 575–584, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Christodoulopoulos, Christos, Sharon Goldwater, and Mark Steedman. 2011. A bayesian mixture model for pos induction using multiple features. In *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing*, pages 638–647, Edinburgh, Scotland, UK., July. Association for Computational Linguistics.
- Clark, Alexander. 2003. Combining distributional and morphological information for part of speech induction. In *Proceedings of the tenth conference on European chapter of the Association for Computational Linguistics - Volume 1*, EACL '03, pages 59–66, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Creutz, Mathias and Krista Lagus. 2005. Inducing the morphological lexicon of a natural language from unannotated text. In *Proceedings of AKRR'05, International and Interdisciplinary Conference on Adaptive Knowledge Representation and Reasoning*, pages 106–113, Espoo, Finland, June.
- Erjavec, Tomaž. 2004. MULTEXT-east version 3: Multilingual morphosyntactic specifications, lexicons and corpora. In *Fourth International Conference on Language Resources and Evaluation, LREC'04*, pages 1535–1538. ELRA.
- Freudenthal, D., J.M. Pine, and F. Gobet. 2005. On the resolution of ambiguities in the extraction of syntactic categories through chunking. *Cognitive Systems Research*, 6(1):17–25.
- Ganchev, Kuzman, João Graça, Jennifer Gillenwater, and Ben Taskar. 2010. Posterior regularization for structured latent variable models. *J. Mach. Learn. Res.*, 99:2001–2049, August.
- Gao, Jianfeng and Mark Johnson. 2008. A comparison of bayesian estimators for unsupervised hidden markov model pos taggers. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, EMNLP '08, pages 344–352, Stroudsburg, PA, USA. Association for Computational Linguistics.

- Globerson, Amir, Gal Chechik, Fernando Pereira, and Naftali Tishby. 2007. Euclidean embedding of co-occurrence data. *J. Mach. Learn. Res.*, 8:2265–2295, December.
- Goldwater, Sharon and Tom Griffiths. 2007. A fully bayesian approach to unsupervised part-of-speech tagging. In *Proceedings of the 45th Annual Meeting of the Association of Computational Linguistics*, pages 744–751, Prague, Czech Republic, June. Association for Computational Linguistics.
- Graff, David, Roni Rosenfeld, and Doug Paul. 1995. Csr-iii text. Linguistic Data Consortium, Philadelphia.
- Haghighi, Aria and Dan Klein. 2006. Prototype-driven learning for sequence models. In *Proceedings of the main conference on Human Language Technology Conference of the North American Chapter of the Association of Computational Linguistics, HLT-NAACL '06*, pages 320–327, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Johnson, Mark. 2007. Why doesn't EM find good HMM POS-taggers? In *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL)*, pages 296–305, Prague, Czech Republic, June. Association for Computational Linguistics.
- Kaufman, L. and P.J. Rousseeuw. 2005. *Finding groups in data: an introduction to cluster analysis*. Wiley series in probability and mathematical statistics. Applied probability and statistics. Wiley.
- Lamar, Michael, Yariv Maron, and Elie Bienenstock. 2010. Latent-descriptor clustering for unsupervised pos induction. In *Proceedings of the 2010 Conference on Empirical Methods in Natural Language Processing, EMNLP '10*, pages 799–809, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Lamar, Michael, Yariv Maron, Mark Johnson, and Elie Bienenstock. 2010. Svd and clustering for unsupervised pos tagging. In *Proceedings of the ACL 2010 Conference Short Papers*, pages 215–219, Uppsala, Sweden, July. Association for Computational Linguistics.
- Lee, Yoong Keok, Aria Haghighi, and Regina Barzilay. 2010. Simple type-level unsupervised pos tagging. In *Proceedings of the 2010 Conference on Empirical Methods in Natural Language Processing, EMNLP '10*, pages 853–861, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Manning, C.D., P. Raghavan, and H. Schütze. 2008. *Introduction to information retrieval*, chapter 17. Cambridge University Press.
- Marcus, Mitchell P., Beatrice Santorini, Mary Ann Marcinkiewicz, and Ann Taylor. 1999. Treebank-3. Linguistic Data Consortium, Philadelphia.
- Maron, Yariv, Michael Lamar, and Elie Bienenstock. 2010. Sphere embedding: An application to part-of-speech induction. In J. Lafferty, C. K. I. Williams, J. Shawe-Taylor, R.S. Zemel, and A. Culotta, editors, *Advances in Neural Information Processing Systems 23*, pages 1567–1575.
- Mintz, T.H. 2003. Frequent frames as a cue for grammatical categories in child directed speech. *Cognition*, 90(1):91–117.
- Ng, A.Y., M.I. Jordan, and Y. Weiss. 2002. On spectral clustering: Analysis and an algorithm. *Advances in neural information processing systems*, 2:849–856.
- Redington, M., N. Crater, and S. Finch. 1998. Distributional information: A powerful cue for acquiring syntactic categories. *Cognitive Science*, 22(4):425–469.
- Rosenberg, A. and J. Hirschberg. 2007. V-measure: A conditional entropy-based external cluster evaluation measure. In *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pages 410–420.
- Roweis, S.T. and L.K. Saul. 2000. Nonlinear dimensionality reduction by locally linear embedding. *Science*, 290(5500):2323.
- Sahlgren, Magnus. 2006. *The Word-Space Model: Using distributional analysis to represent syntagmatic and paradigmatic relations between words in high-dimensional vector spaces*. Ph.D. thesis, Stockholm University.
- Sak, H., T. Güngör, and M. Saraçlar. 2008. Turkish language resources: Morphological parser, morphological disambiguator and web corpus. *Advances in natural language processing*, pages 417–427.
- Schütze, H. and J. Pedersen. 1993. A Vector Model for syntagmatic and paradigmatic relatedness. In *Proceedings of the 9th Annual Conference of the University of Waterloo Centre for the New OED and Text Research*, Oxford, England.
- Schütze, Hinrich. 1995. Distributional part-of-speech tagging. In *Proceedings of the seventh conference on European chapter of the Association for Computational Linguistics, EACL '95*, pages 141–148, San Francisco, CA, USA. Morgan Kaufmann Publishers Inc.
- Stolcke, Andreas. 2002. Srilm-an extensible language modeling toolkit. In *Proceedings International Conference on Spoken Language Processing*, pages 257–286, November.

- Tenenbaum, J.B., V. Silva, and J.C. Langford. 2000. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319.
- Toutanova, Kristina, Dan Klein, Christopher D. Manning, and Yoram Singer. 2003. Feature-rich part-of-speech tagging with a cyclic dependency network. In *Proceedings of the 2003 Conference of the North American Chapter of the Association for Computational Linguistics on Human Language Technology - Volume 1*, NAACL '03, pages 173–180, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Turney, Peter D. and Patrick Pantel. 2010. From frequency to meaning: Vector space models of semantics. *J. Artif. Intell. Res. (JAIR)*, 37:141–188.
- Yatbaz, Mehmet Ali, Enis Sert, and Deniz Yuret. 2012. Learning syntactic categories using paradigmatic representations of word context. In *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pages 940–951, Jeju Island, Korea, July. Association for Computational Linguistics.
- Yatbaz, Mehmet Ali and Deniz Yuret. 2010. Unsupervised part of speech tagging using unambiguous substitutes from a statistical language model. In *Coling 2010: Posters*, pages 1391–1398, Beijing, China, August. Coling 2010 Organizing Committee.
- Yuret, Deniz. 2012. Fastsubs: An efficient and exact procedure for finding the most likely lexical substitutes based on an n-gram language model. *To appear in IEEE Signal Processing Letters*.

—

—

—

—

—

—